

Foundation models

For autonomous driving

Vincent Vanhoucke
Distinguished Engineer
Waymo



OUR MISSION –

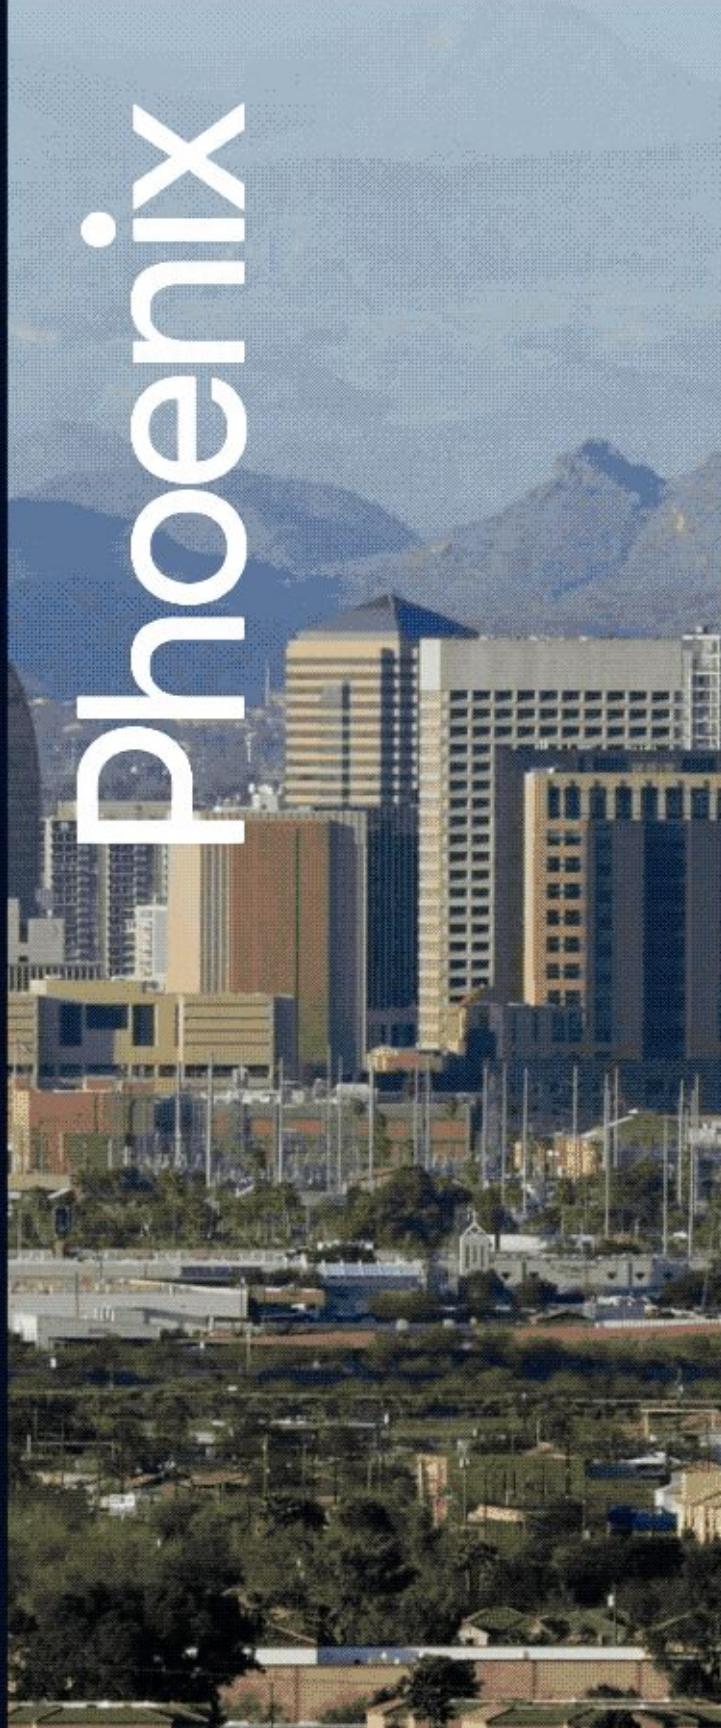
Be the world's
most trusted driver.



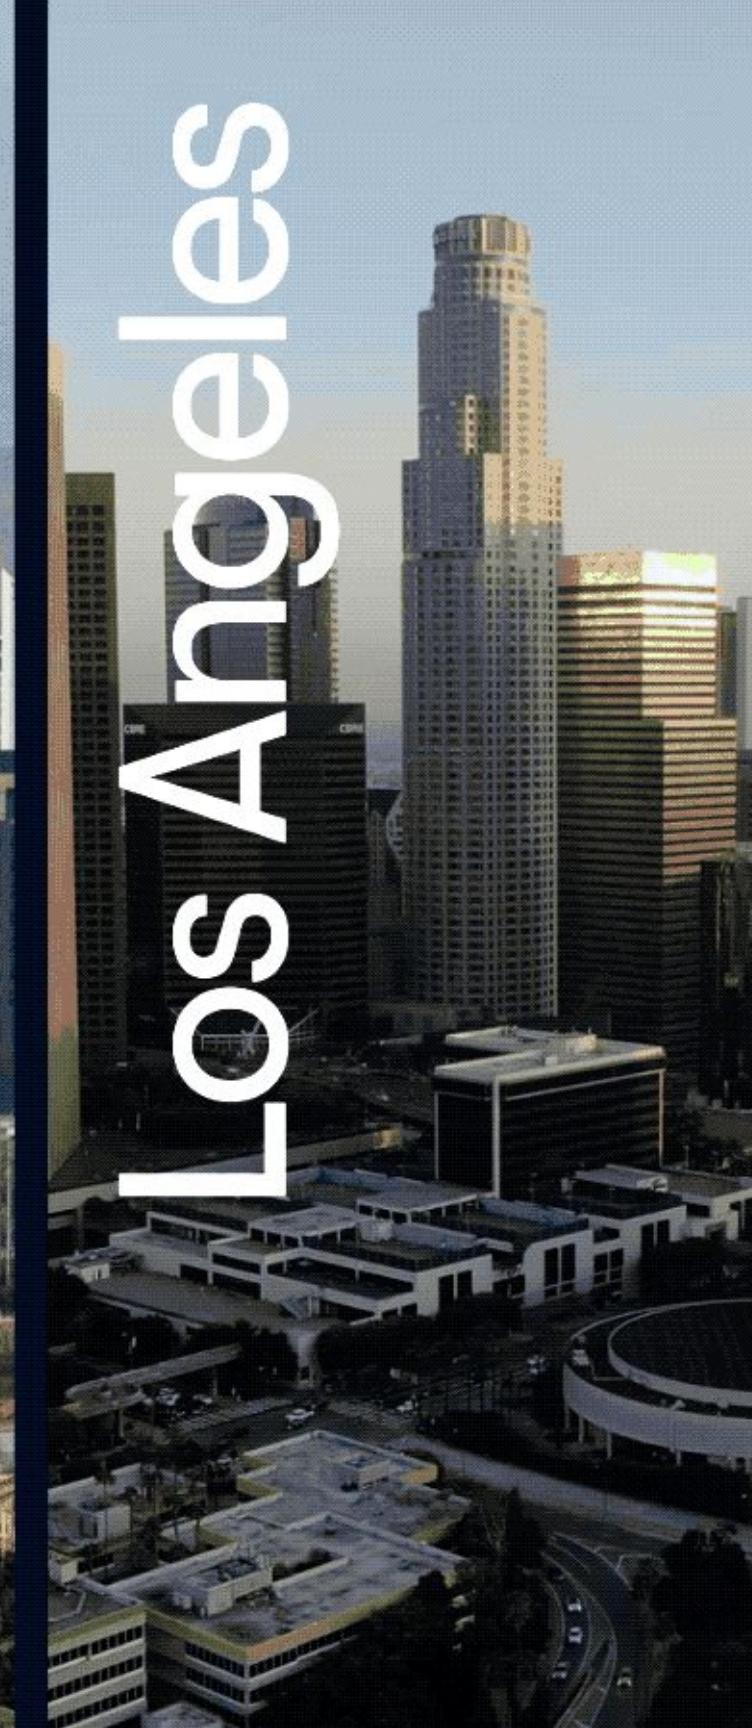
San Francisco



Phoenix



Los Angeles



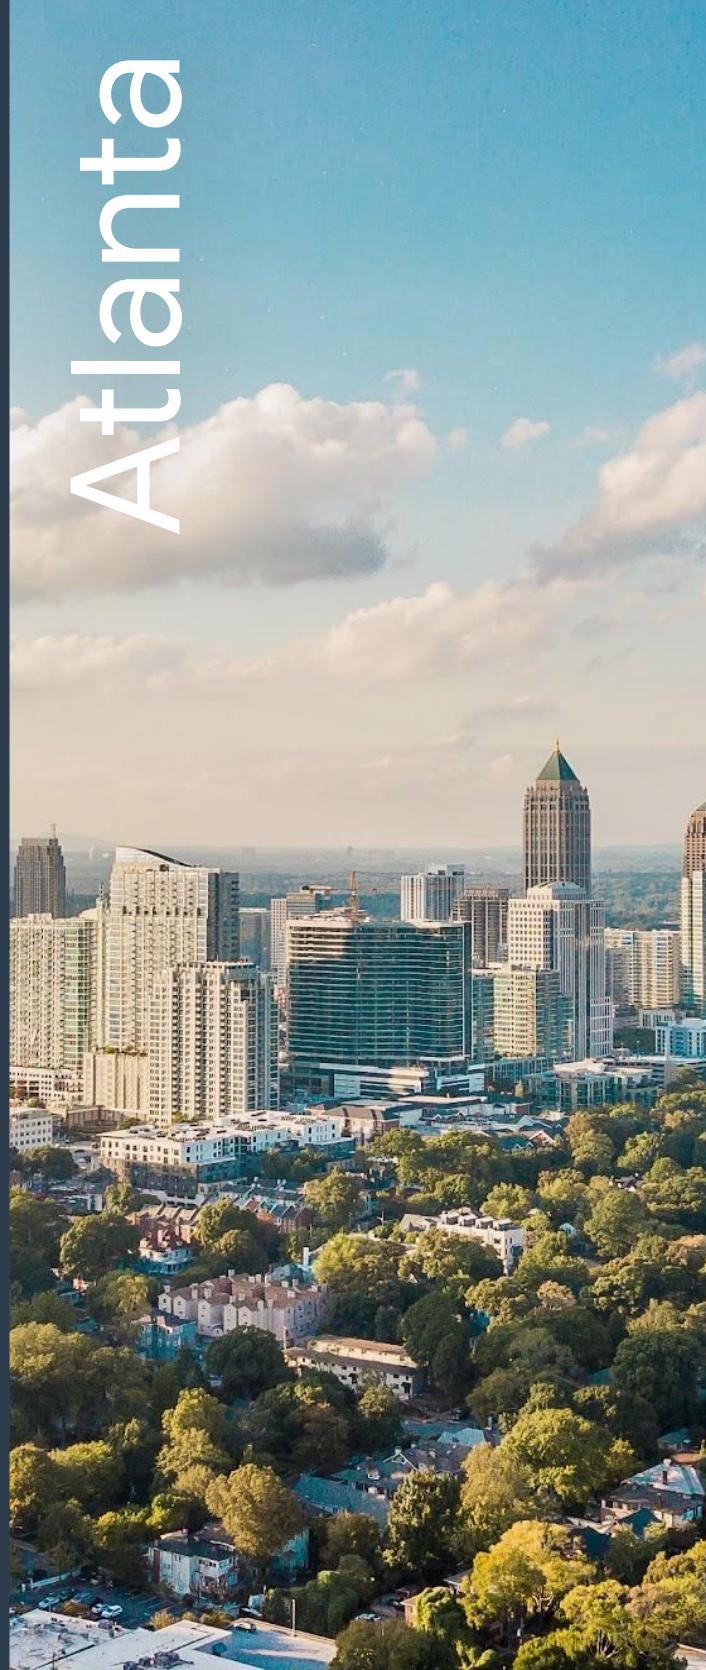
Austin



Up next

WAYMO ONE x Uber

Atlanta



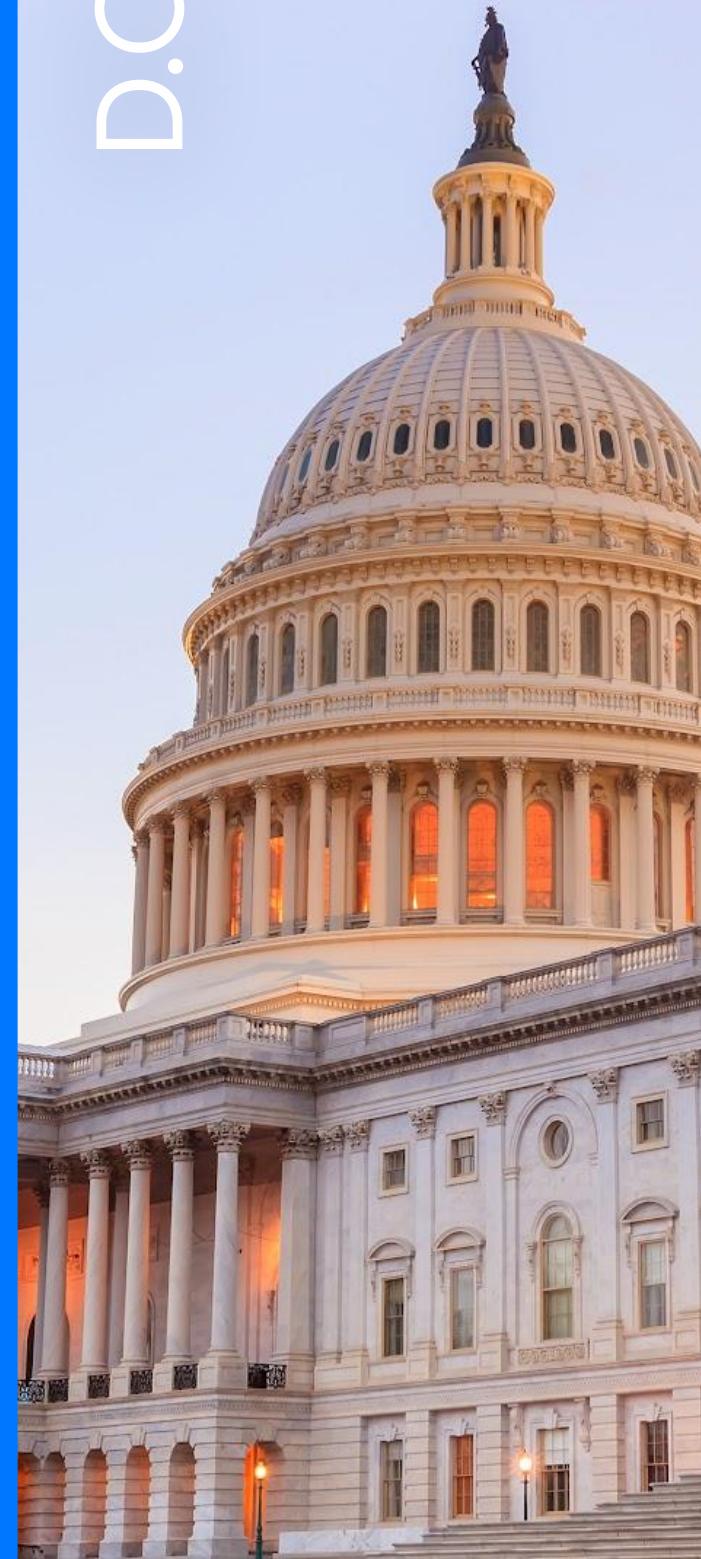
WAYMO ONE

Miami

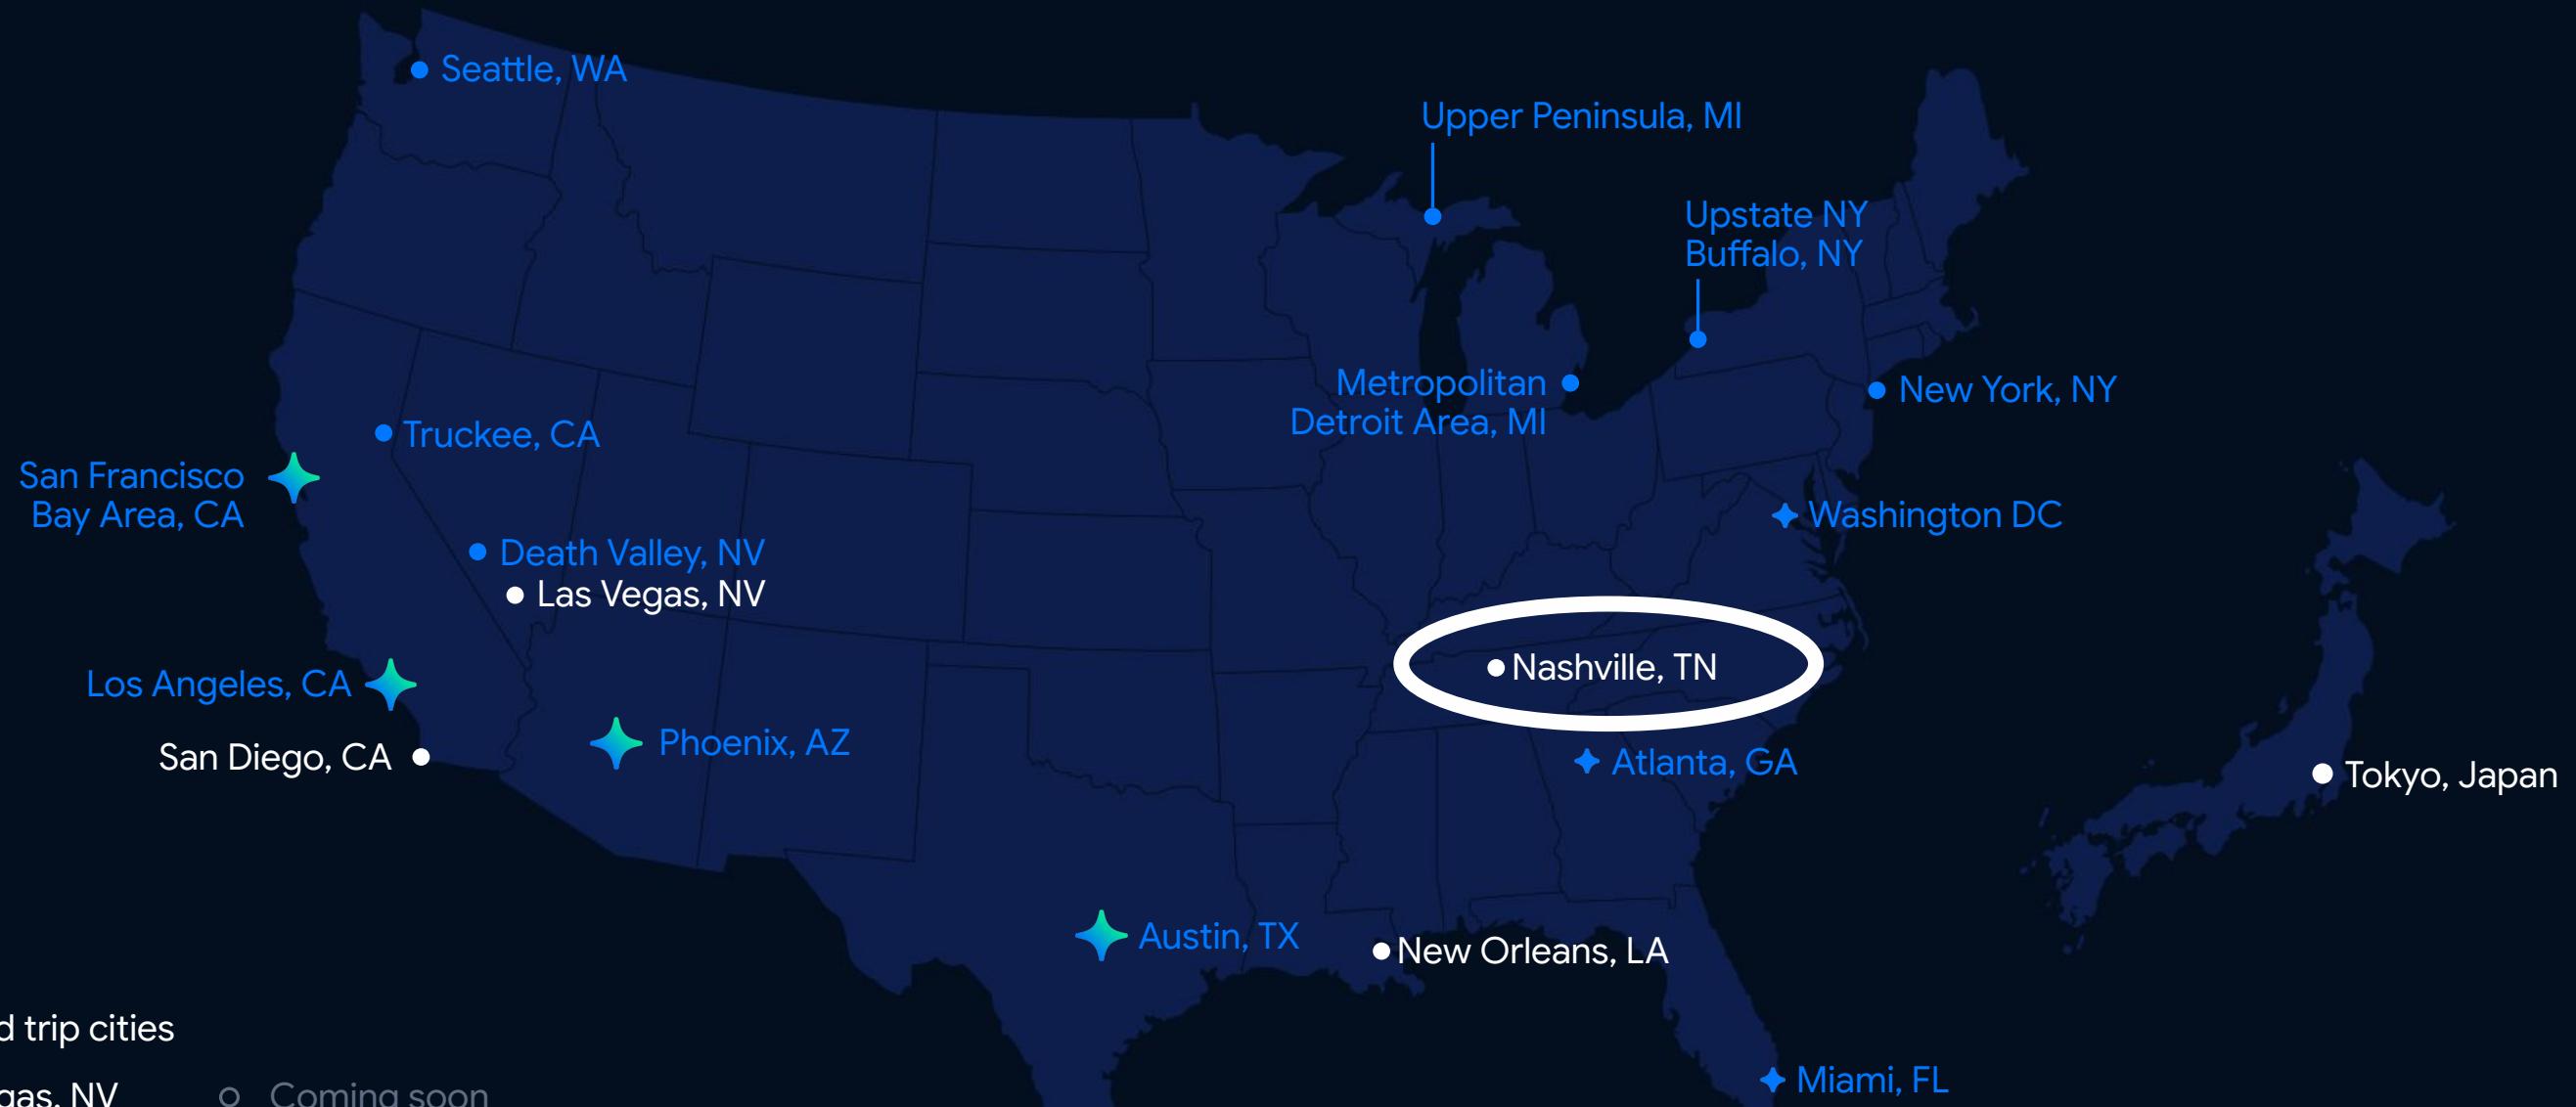


WAYMO ONE

D.C.



◆ Waymo operating cities ♦ Operation coming soon ● Waymo tested cities



2025 Road trip cities

- | | |
|-------------------|---------------|
| ● Las Vegas, NV | ○ Coming soon |
| ● San, Diego, CA | ○ Coming soon |
| ● New Orleans, LA | ○ Coming soon |
| ● Nashville, TN | ○ Coming soon |
| ○ Coming soon | ○ Coming soon |

Road trip[★]

WAYMO ONE

Now
serving over

250,000

paid trips
per week



Protecting Vulnerable Road Users

Compared to human drivers over **56.7 million** miles in our operational cities, Waymo Driver had:



92%

Fewer crashes with
injuries to pedestrians



82%

Fewer crashes with
injuries to cyclists



82%

Fewer crashes with
injuries to motorcyclists

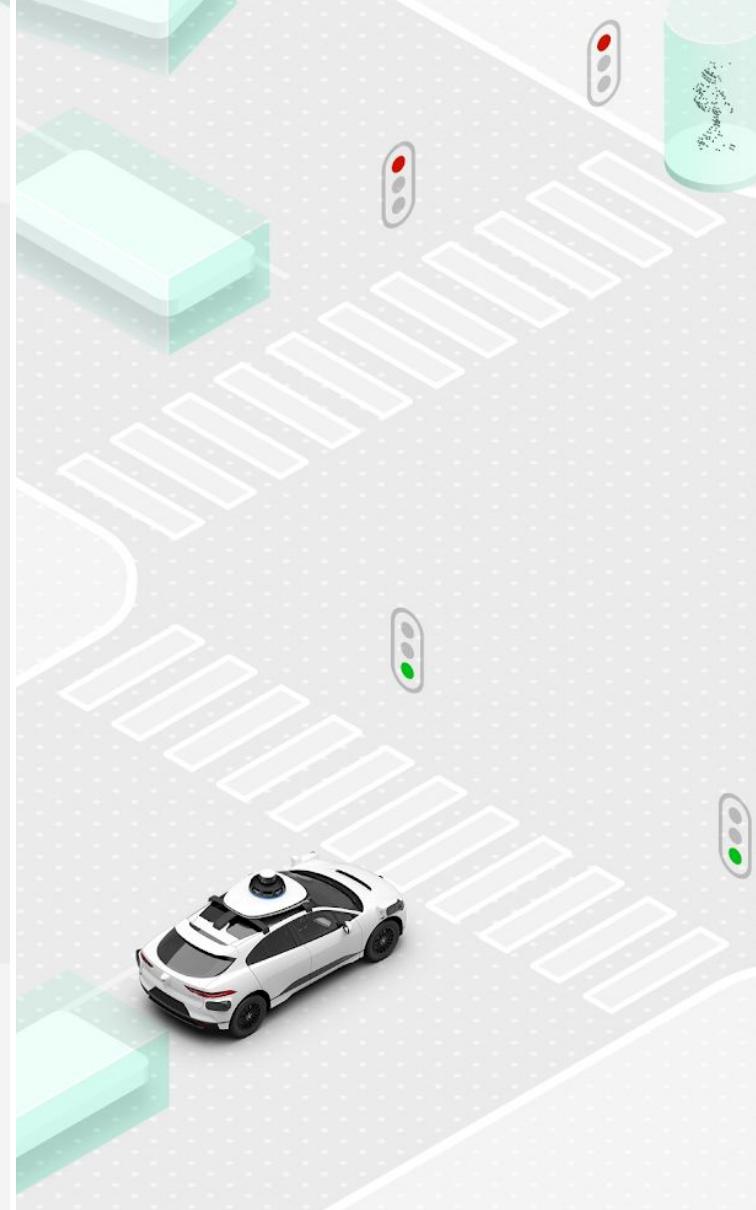
STEP 1

Where am I?



STEP 2

What's
around me?



STEP 3

What will
happen next?



STEP 4

What
should I do?





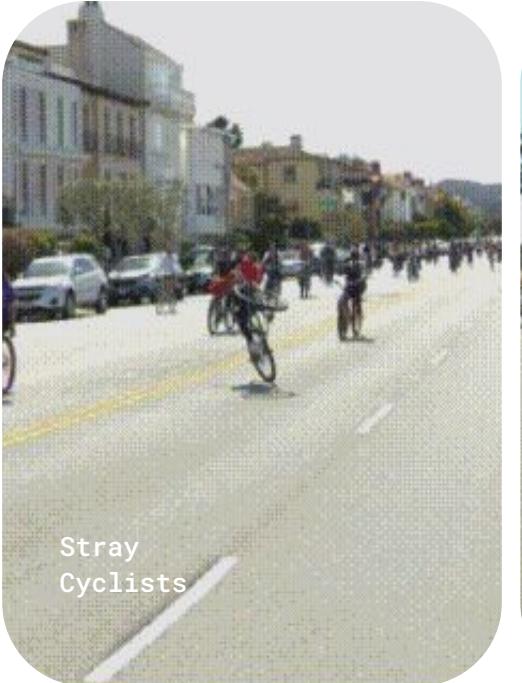
Long tail of driving scenarios presents significant challenges

Extreme Weather



Falling Trees

Unusual Behaviors

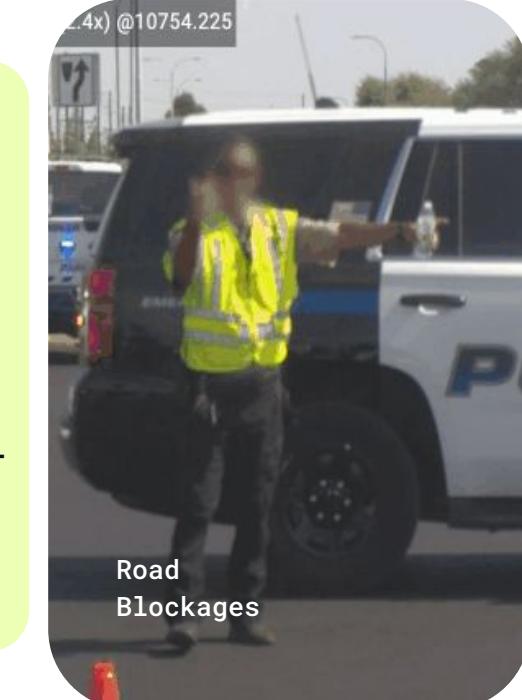


Stray Cyclists

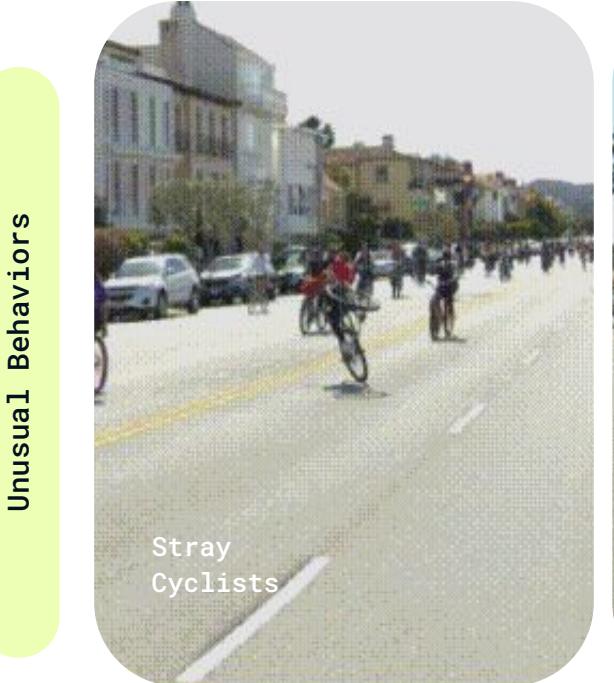


Flooding

Unique Interactions



Road Blockages



Toddlers on the Loose

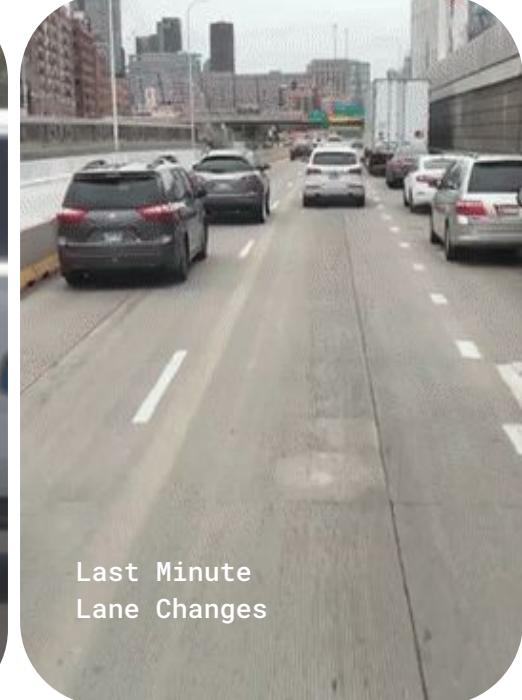


Falling Skateboarders

Foreign Objects on the Road



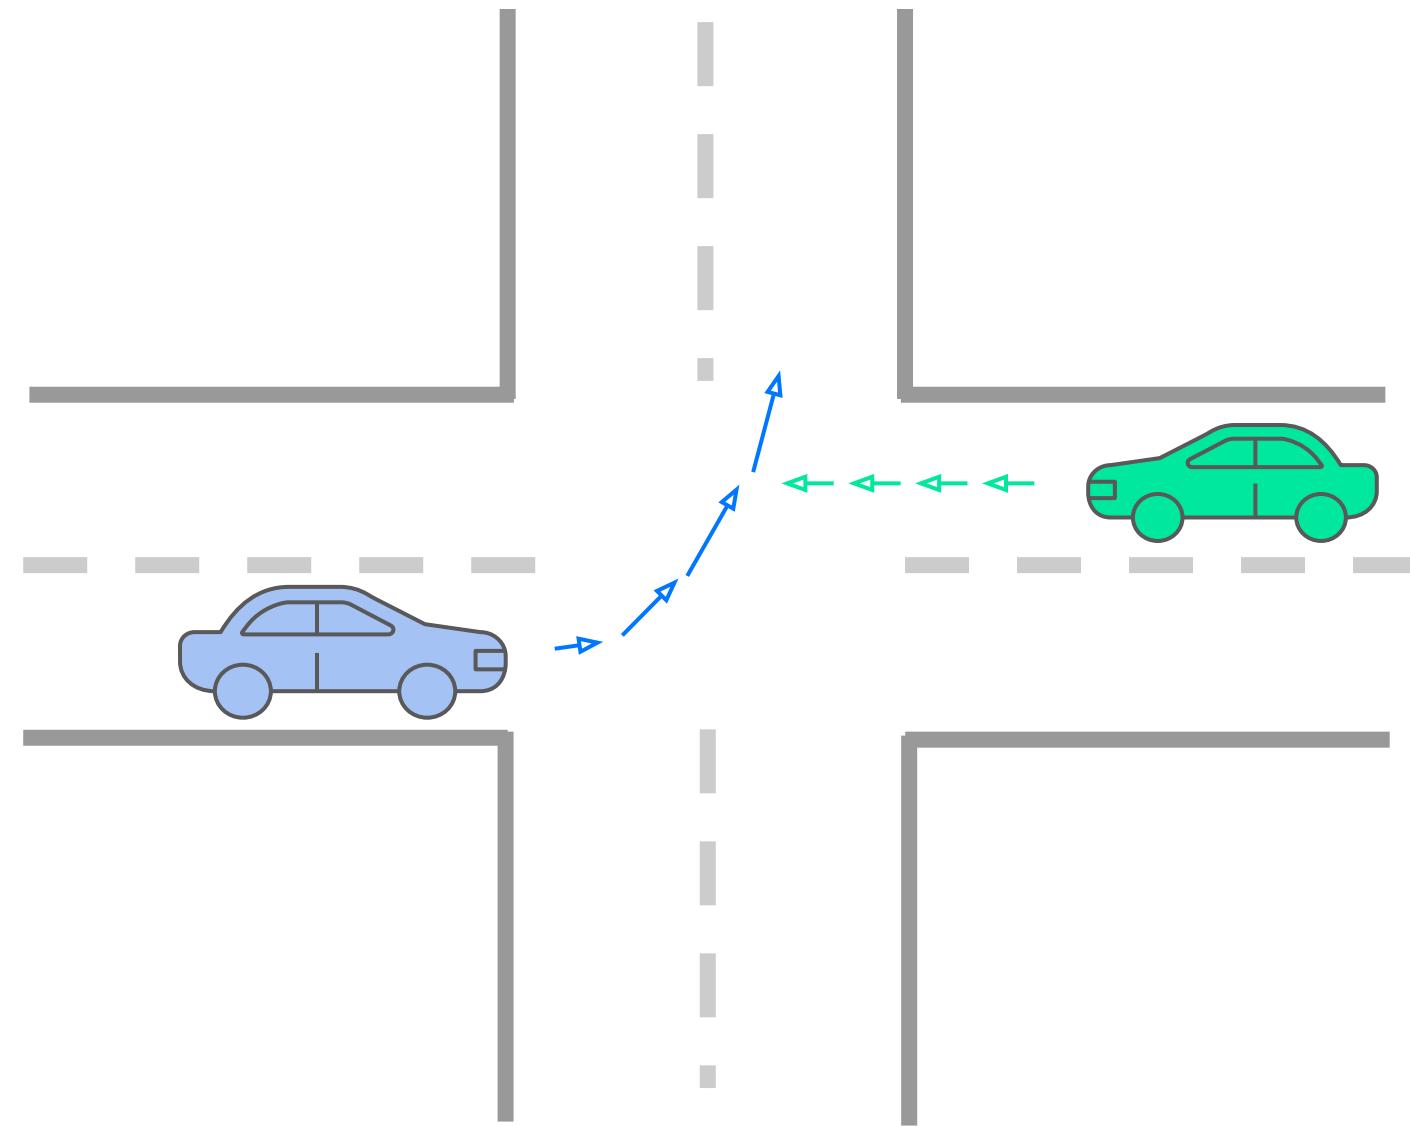
A BBQ Grill Falling Off Truck



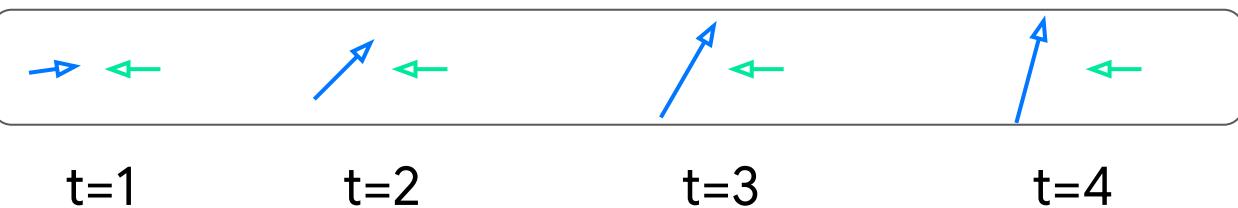
Last Minute Lane Changes

Driving as a Conversation

- Trajectories as sentences in a **new language**
- Vocabulary consists of **state / motion words** (vectors)
- Like language, trajectories have **local continuity and global context**
- Model architecture is **similar to** that of a **LLM**



Motion token sequence:



MotionLM: Multi-Agent Motion Forecasting as Language Modeling

Ari Seff, Brian Cera, Dian Chen, Mason Ng, Aurick Zhou, Nigamaa Nayakanti, Khaled S. Refaat, Rami Al-Rfou, Benjamin Sapp

ICCV 2023



Examples

MARGINAL PREDICTION

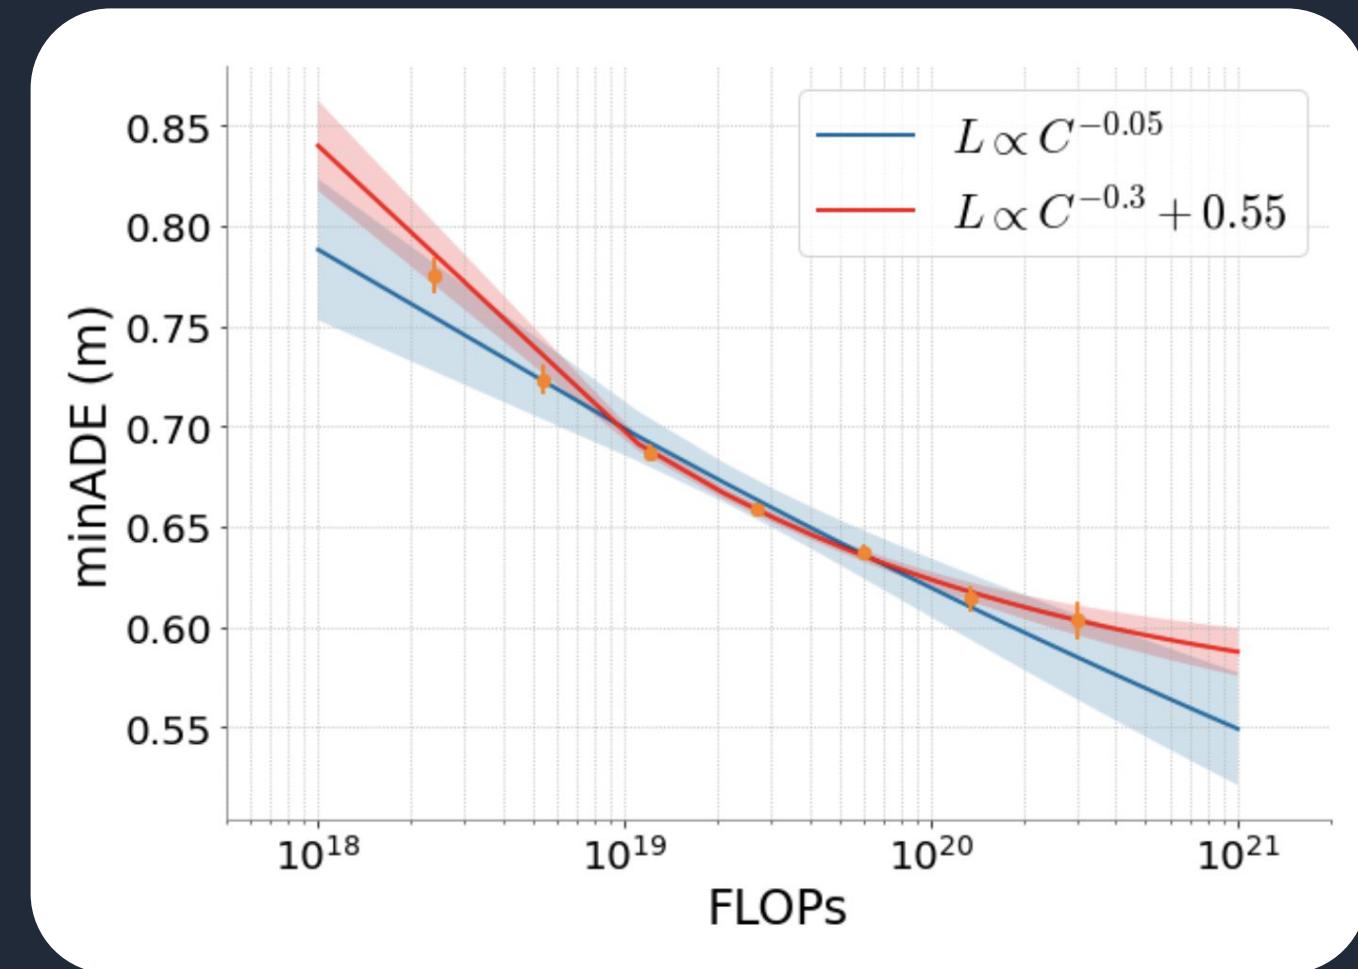
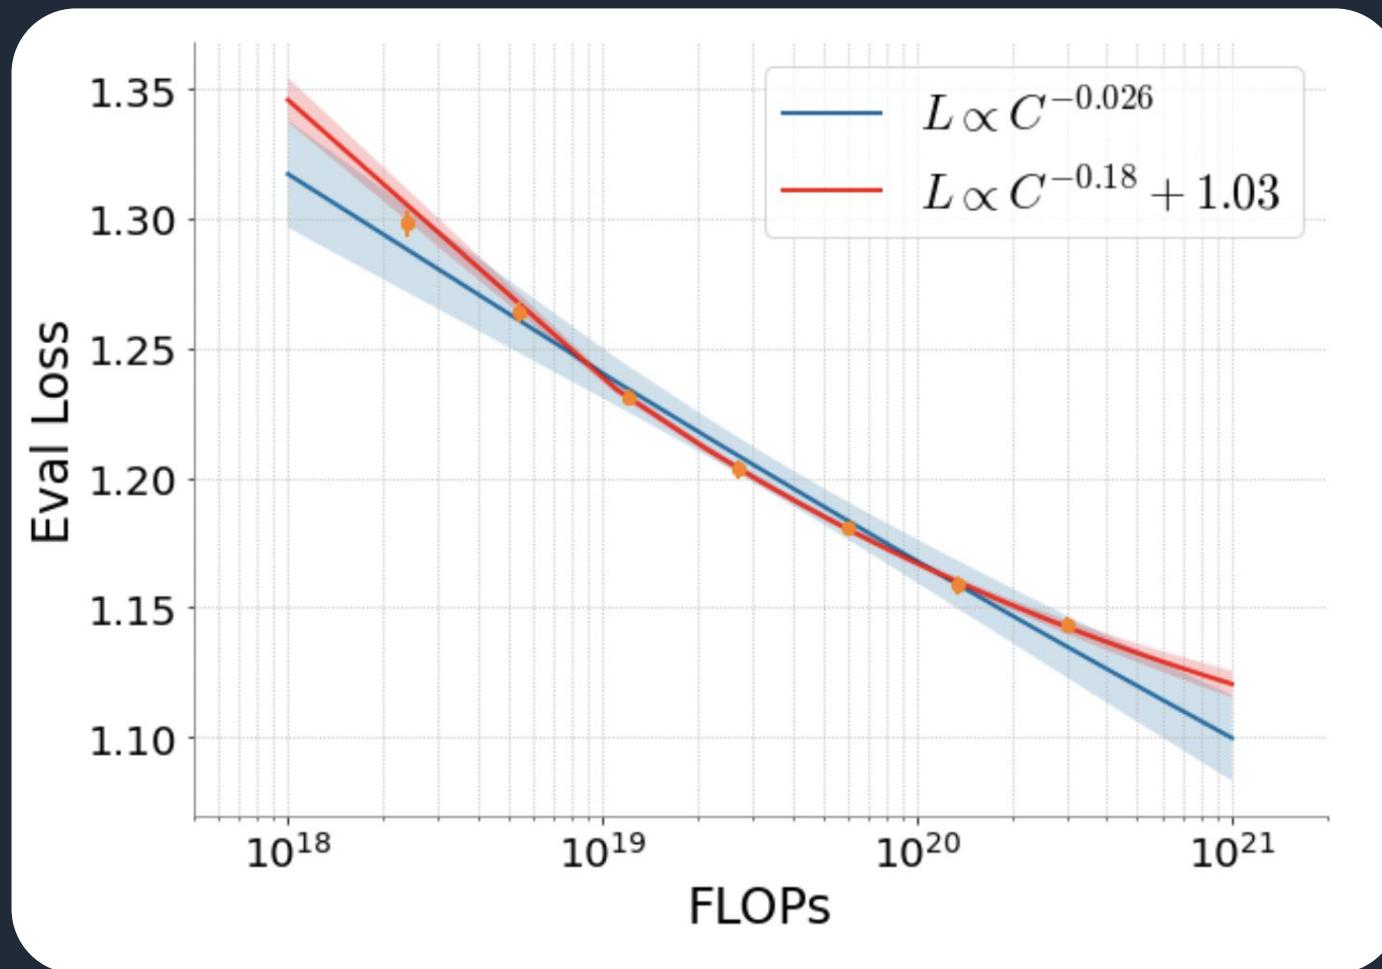


CAUSAL JOINT “DIALOGUE”



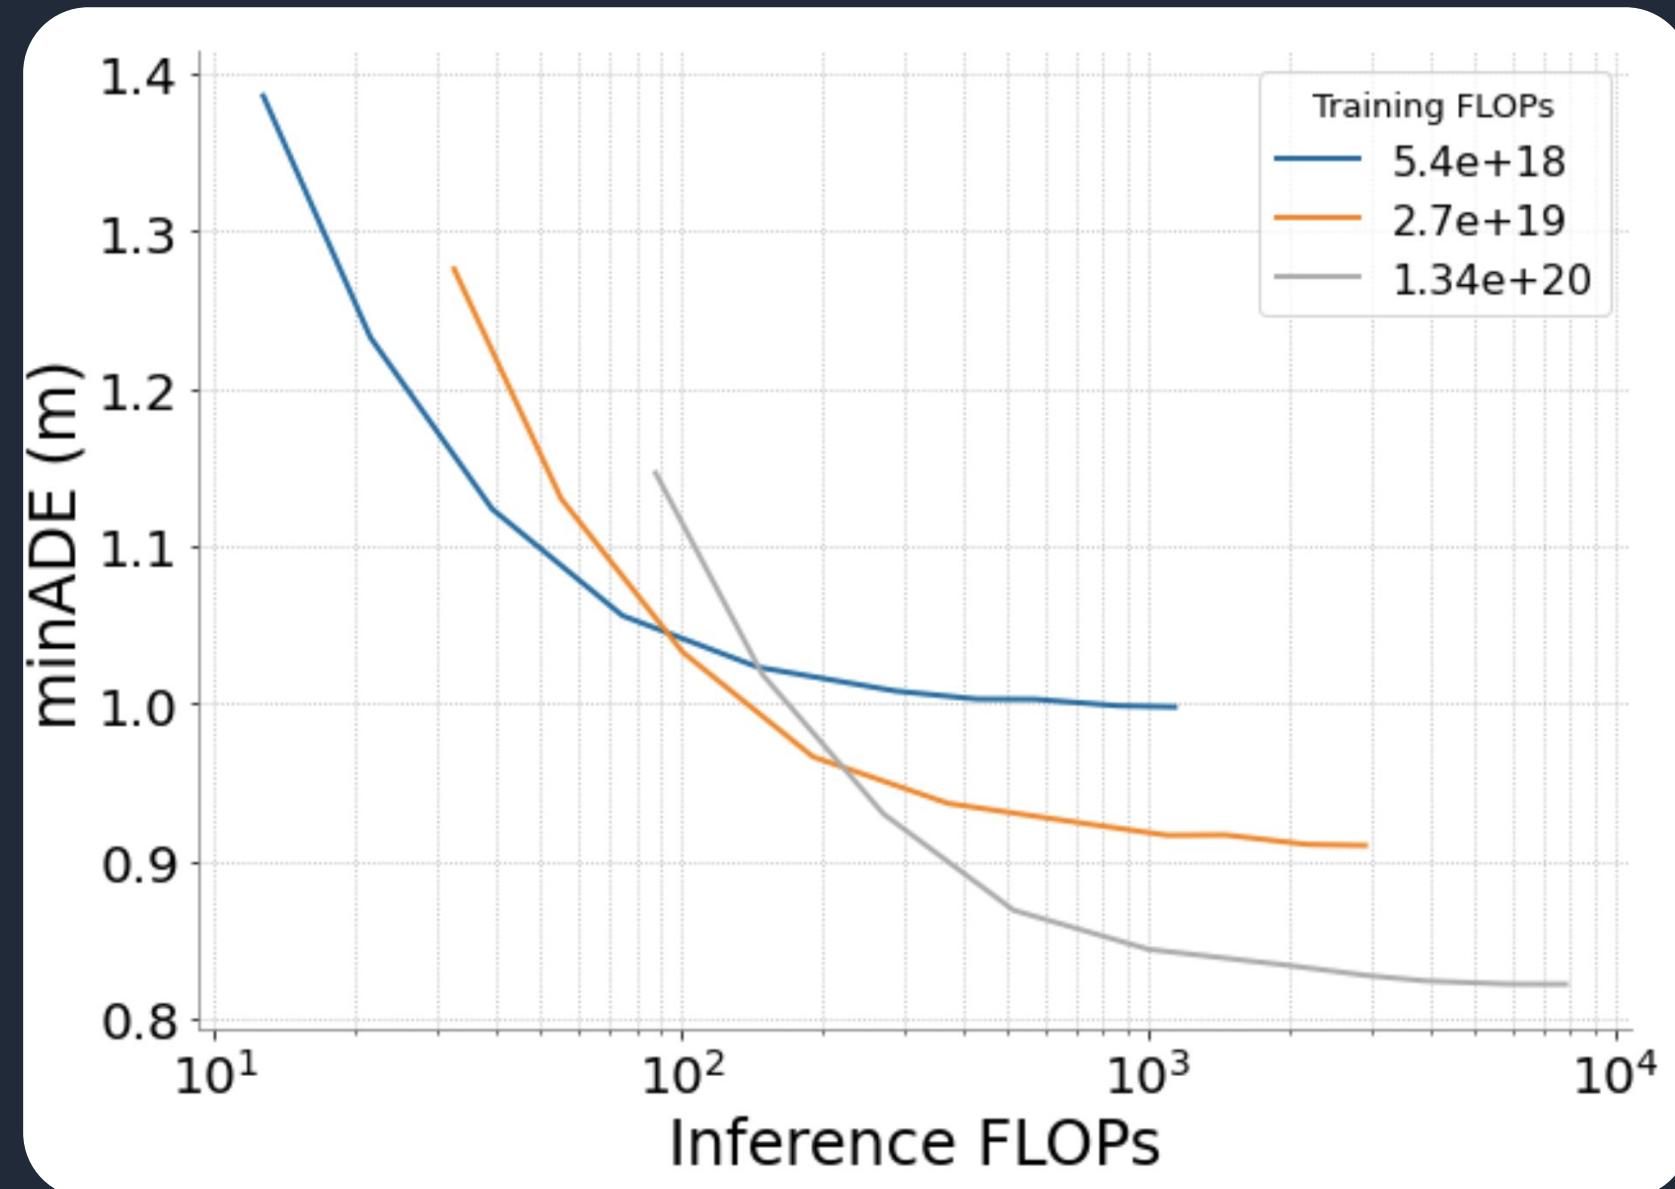
Scaling Laws

Once you have a good architecture, performance scales with model size and data



Inference Scaling Laws

Performance scales with the amount of compute used at runtime



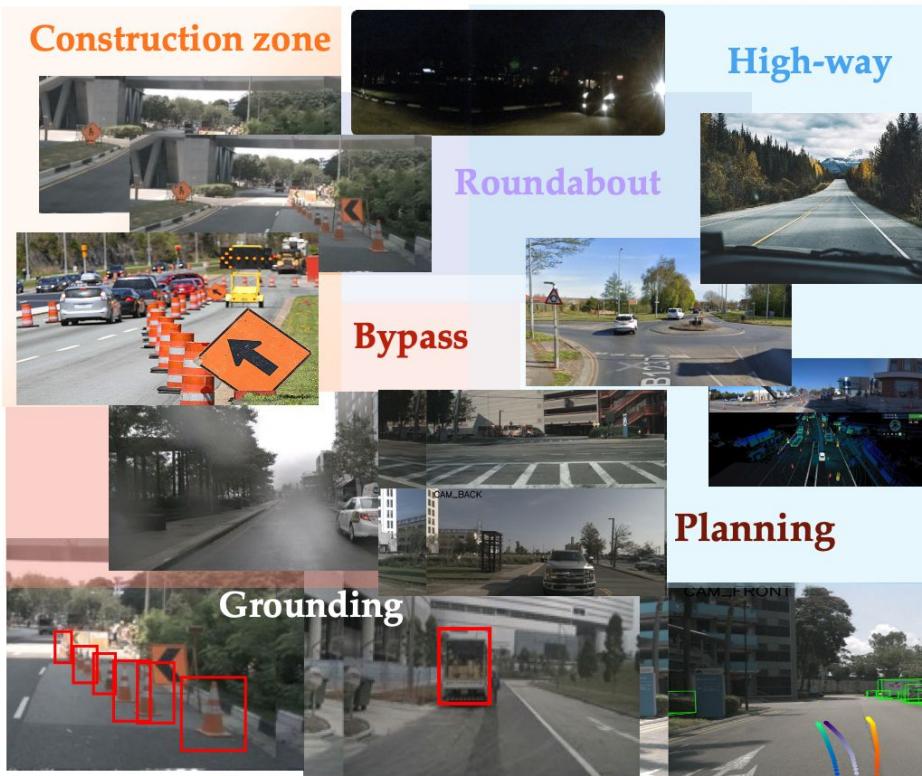
For more on motion scaling laws,
don't miss **Ben Sapp's** talk
tomorrow at the Workshop on
Distillation of Foundation Models
for Autonomous Driving!

(NEW!) Now on ArXiv

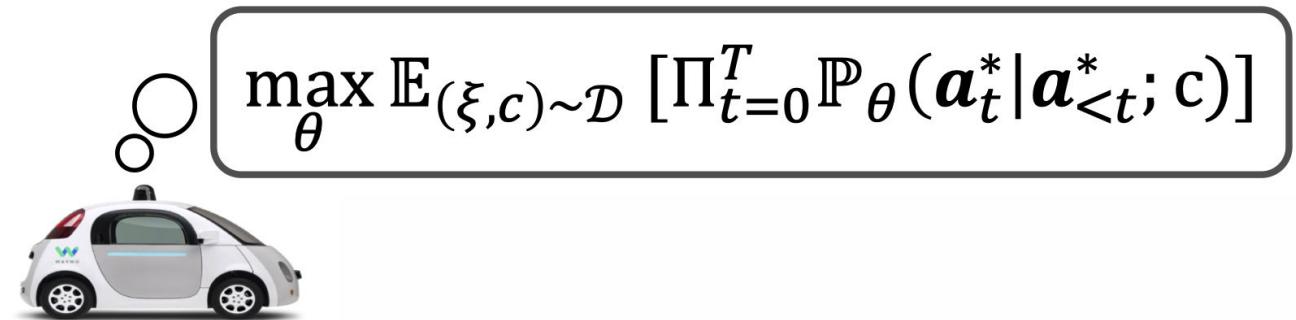


Post-training Preference Alignment

Large scale driving demonstration data



Let's remember what the expert did and copy them!



Misalignment: by optimizing an incomplete or mis-specified objective, these models lead to undesirable behaviors at best and safety hazards at worst!

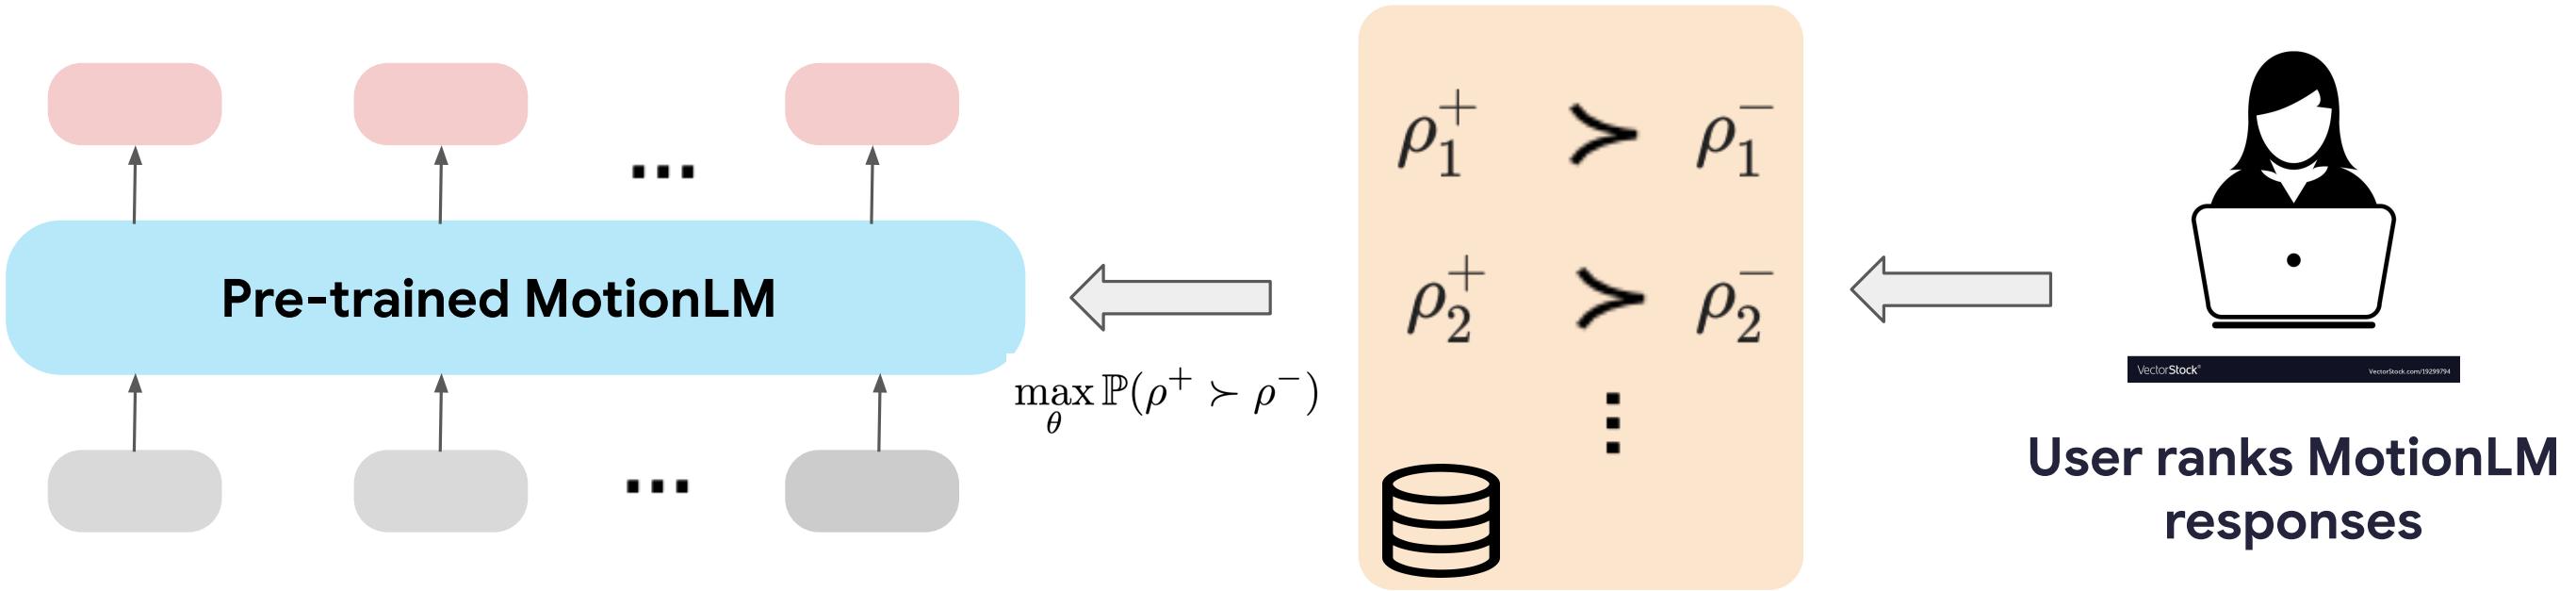
[Direct Post-Training Preference Alignment for Multi-Agent Motion Generation Model](#)

Thomas (Ran) Tian, Kratarth Goel

ICLR 2025, Spotlight



Post-training Preference Alignment



Reconcile the disparity between the next-token prediction objective and human preferences.

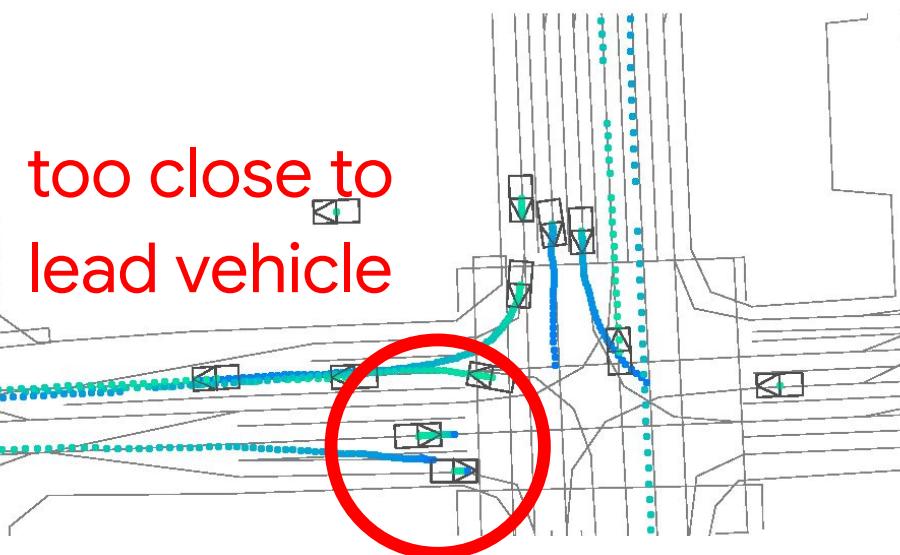
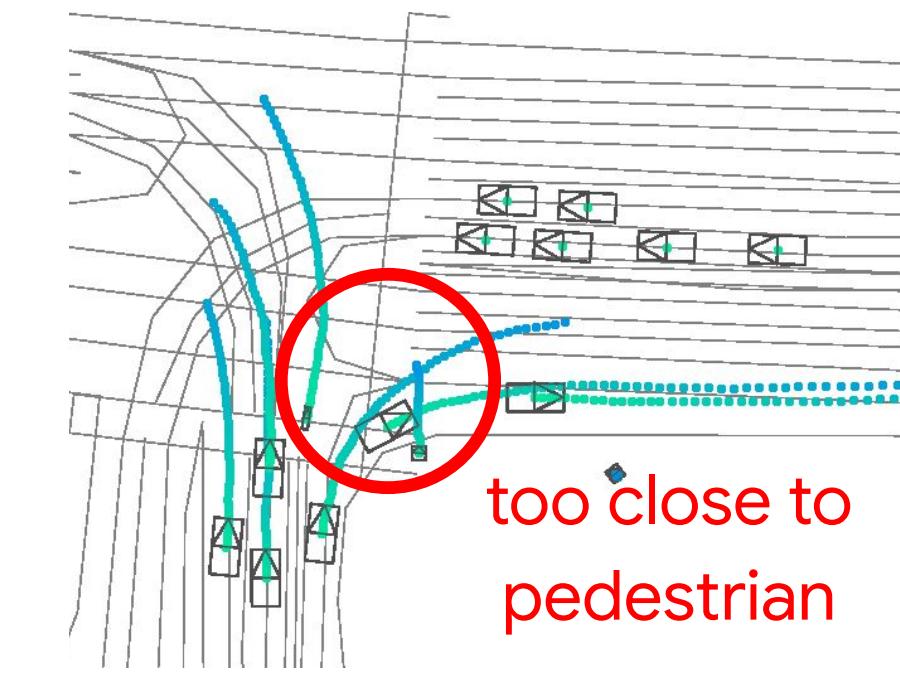
Direct Post-Training Preference Alignment for Multi-Agent Motion Generation Model

Using Implicit Feedback from Pre-training Demonstrations

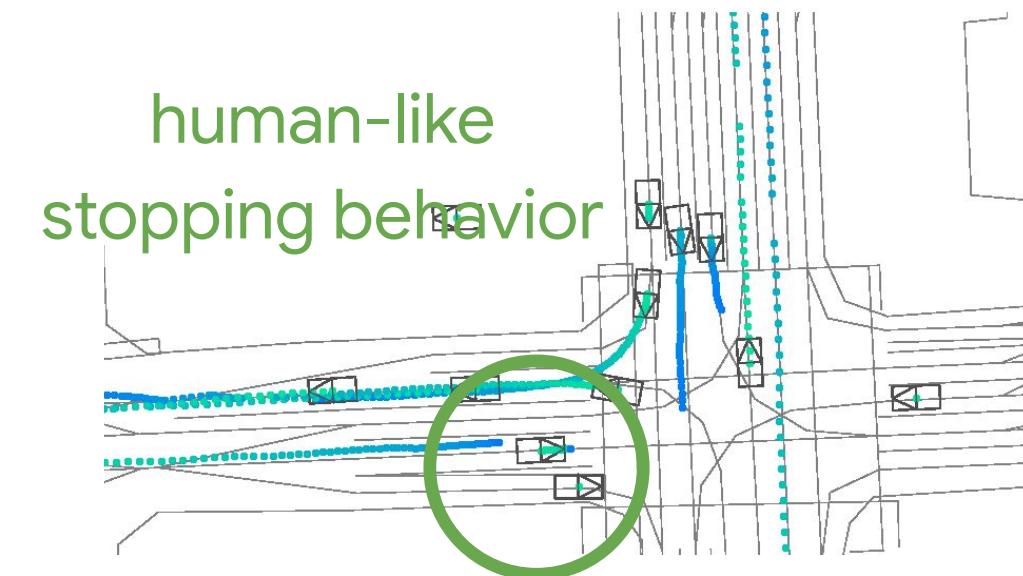
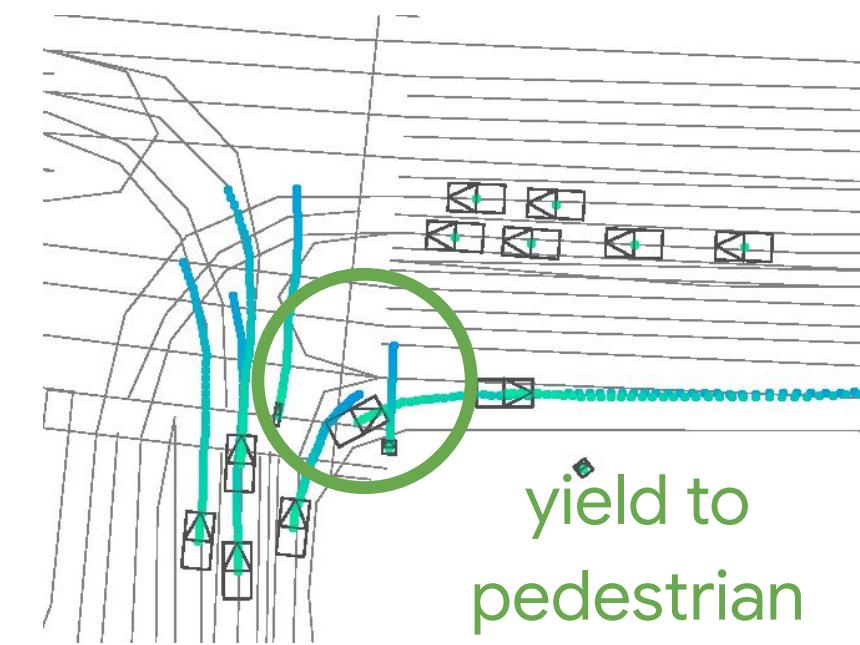
Thomas (Ran) Tian, Kratarth Goel

ICLR 2025, Spotlight

Pre-trained traffic simulation model



After post-training alignment



Leveraging language understanding and general reasoning capabilities of LLM/VLMs



**It's Wednesday at 4pm. Can I park at this spot right now?
Tell me in 1 line.**

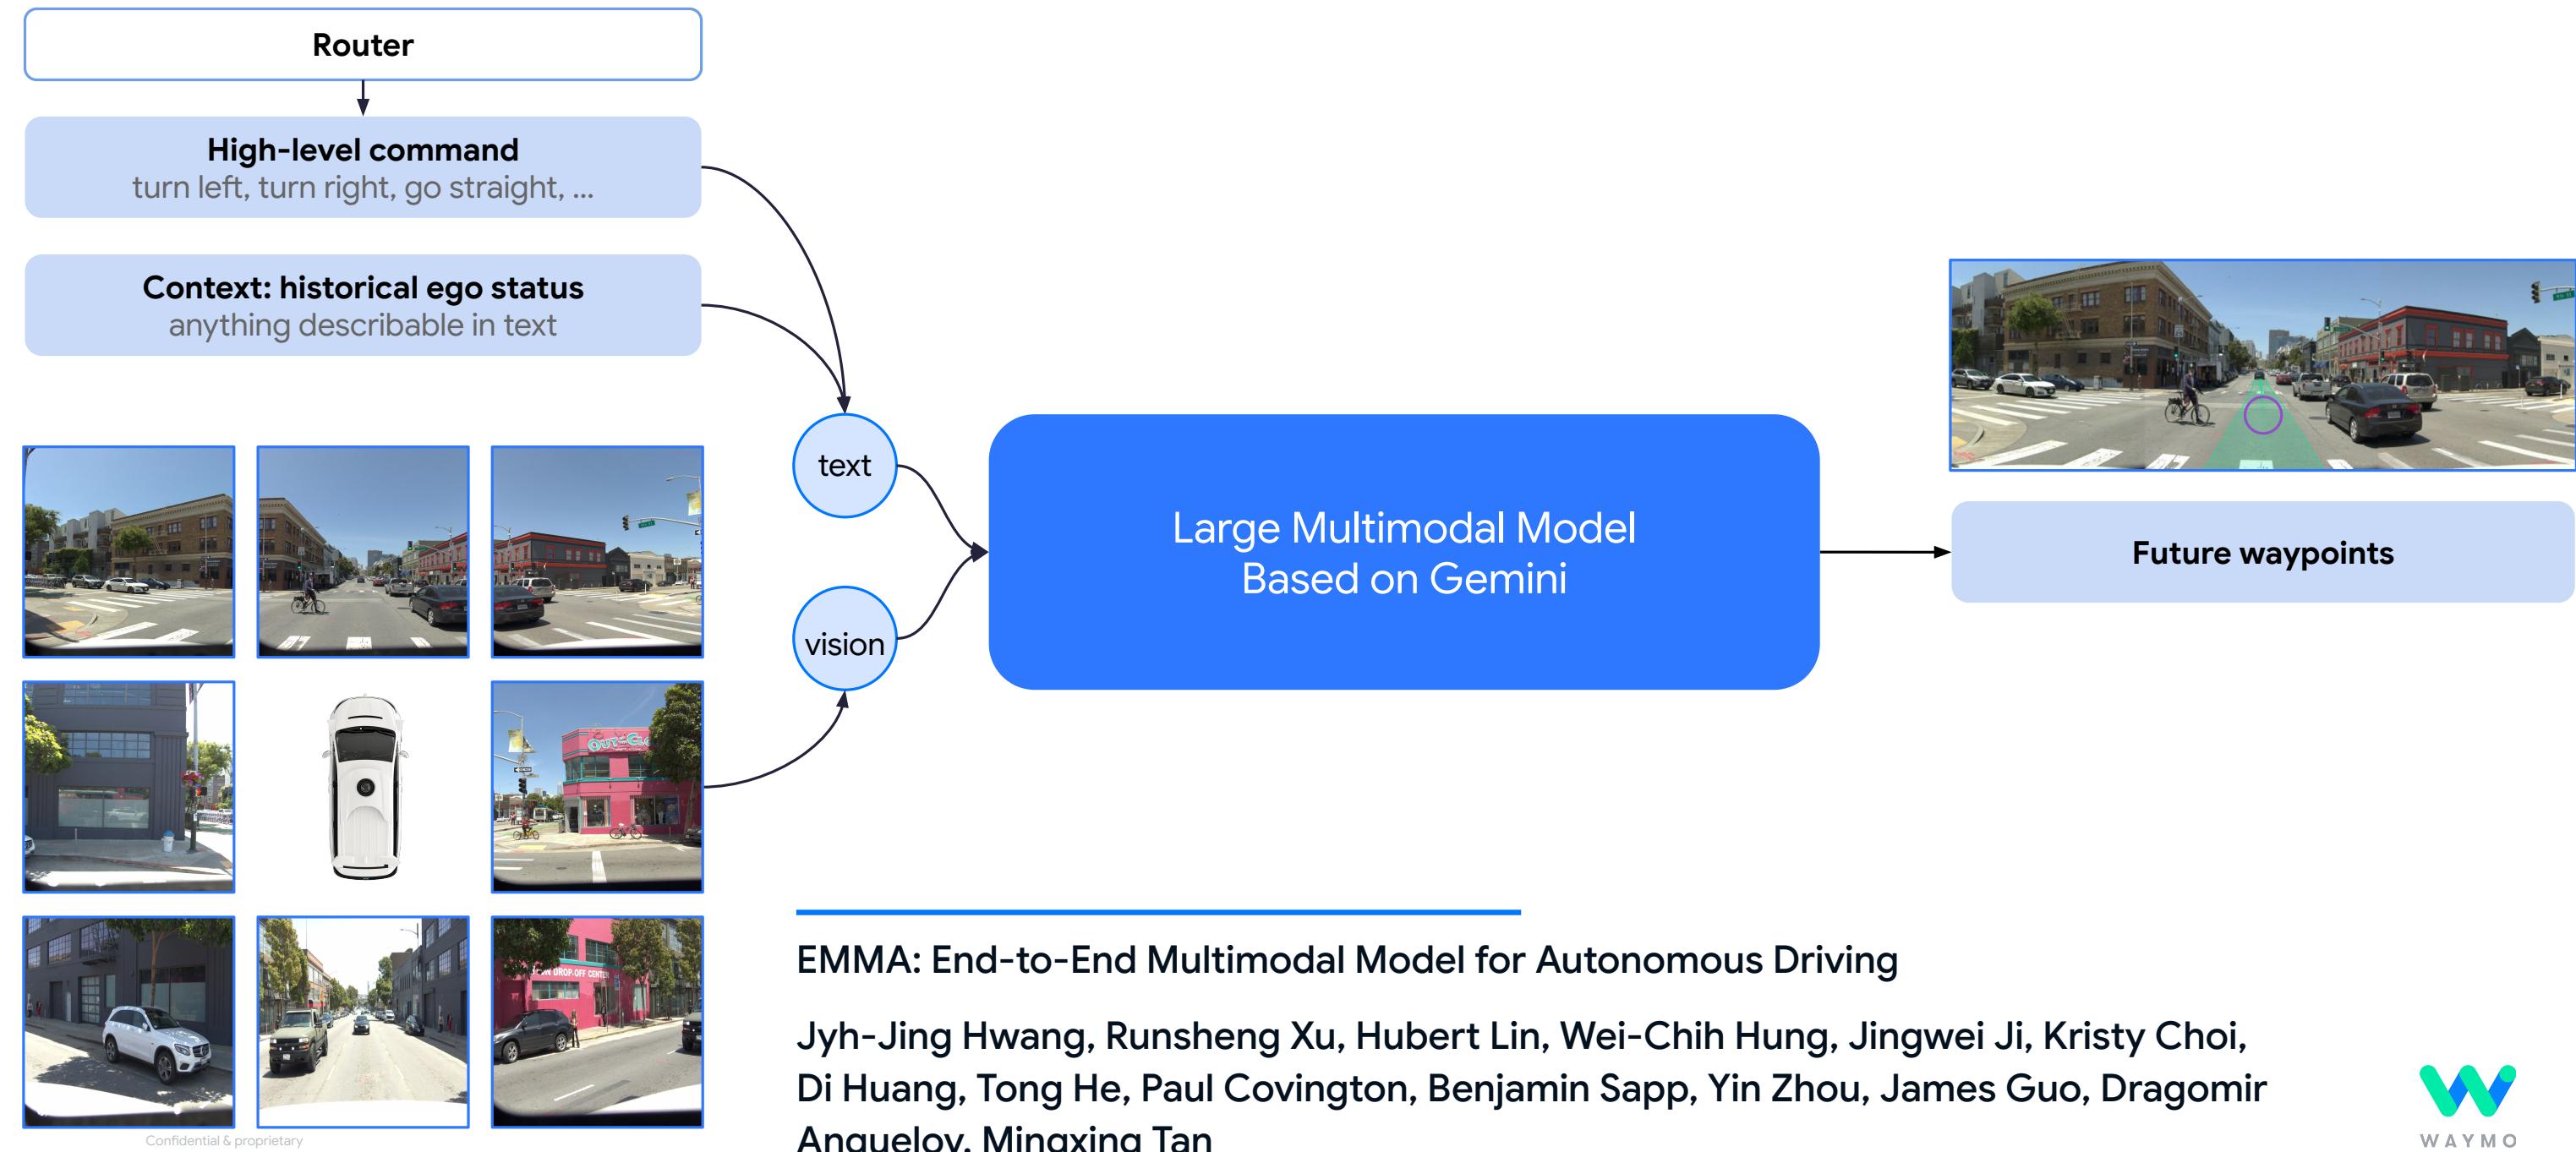
VLM: Yes, you can park for up to 1 hour starting at 4pm.



I am driving on the street, this is what I see from the front window. Tell me what's happening and what I should pay special attention to. And tell me what I should do and specify the reasons.

VLM: There's been a serious accident! A car has flipped onto its roof and is blocking the streetcar tracks. Emergency responders are on the scene. Pay close attention to the emergency personnel and any instructions they may give. You should slow down significantly, proceed with extreme caution, and be prepared to stop completely if necessary. The most important thing is to avoid hindering the emergency response and to ensure your own safety and the safety of others around you. If possible, consider an alternate route to avoid the area entirely.

Multimodal Models for AV motion prediction



Competitive Academic Benchmark Results

WAYMO OPEN MOTION DATASET

Method	ADE 1s	ADE 3s	ADE 5s	ADE 8s
MotionLM [Seff et al., 2023]	0.045	0.251	0.694	1.766
Wayformer* [Nayakanti et al., 2023]	0.044	0.250	0.640	1.517
EMMA	0.032	0.248	0.681	1.718
EMMA+	0.030	0.225	0.610	1.553

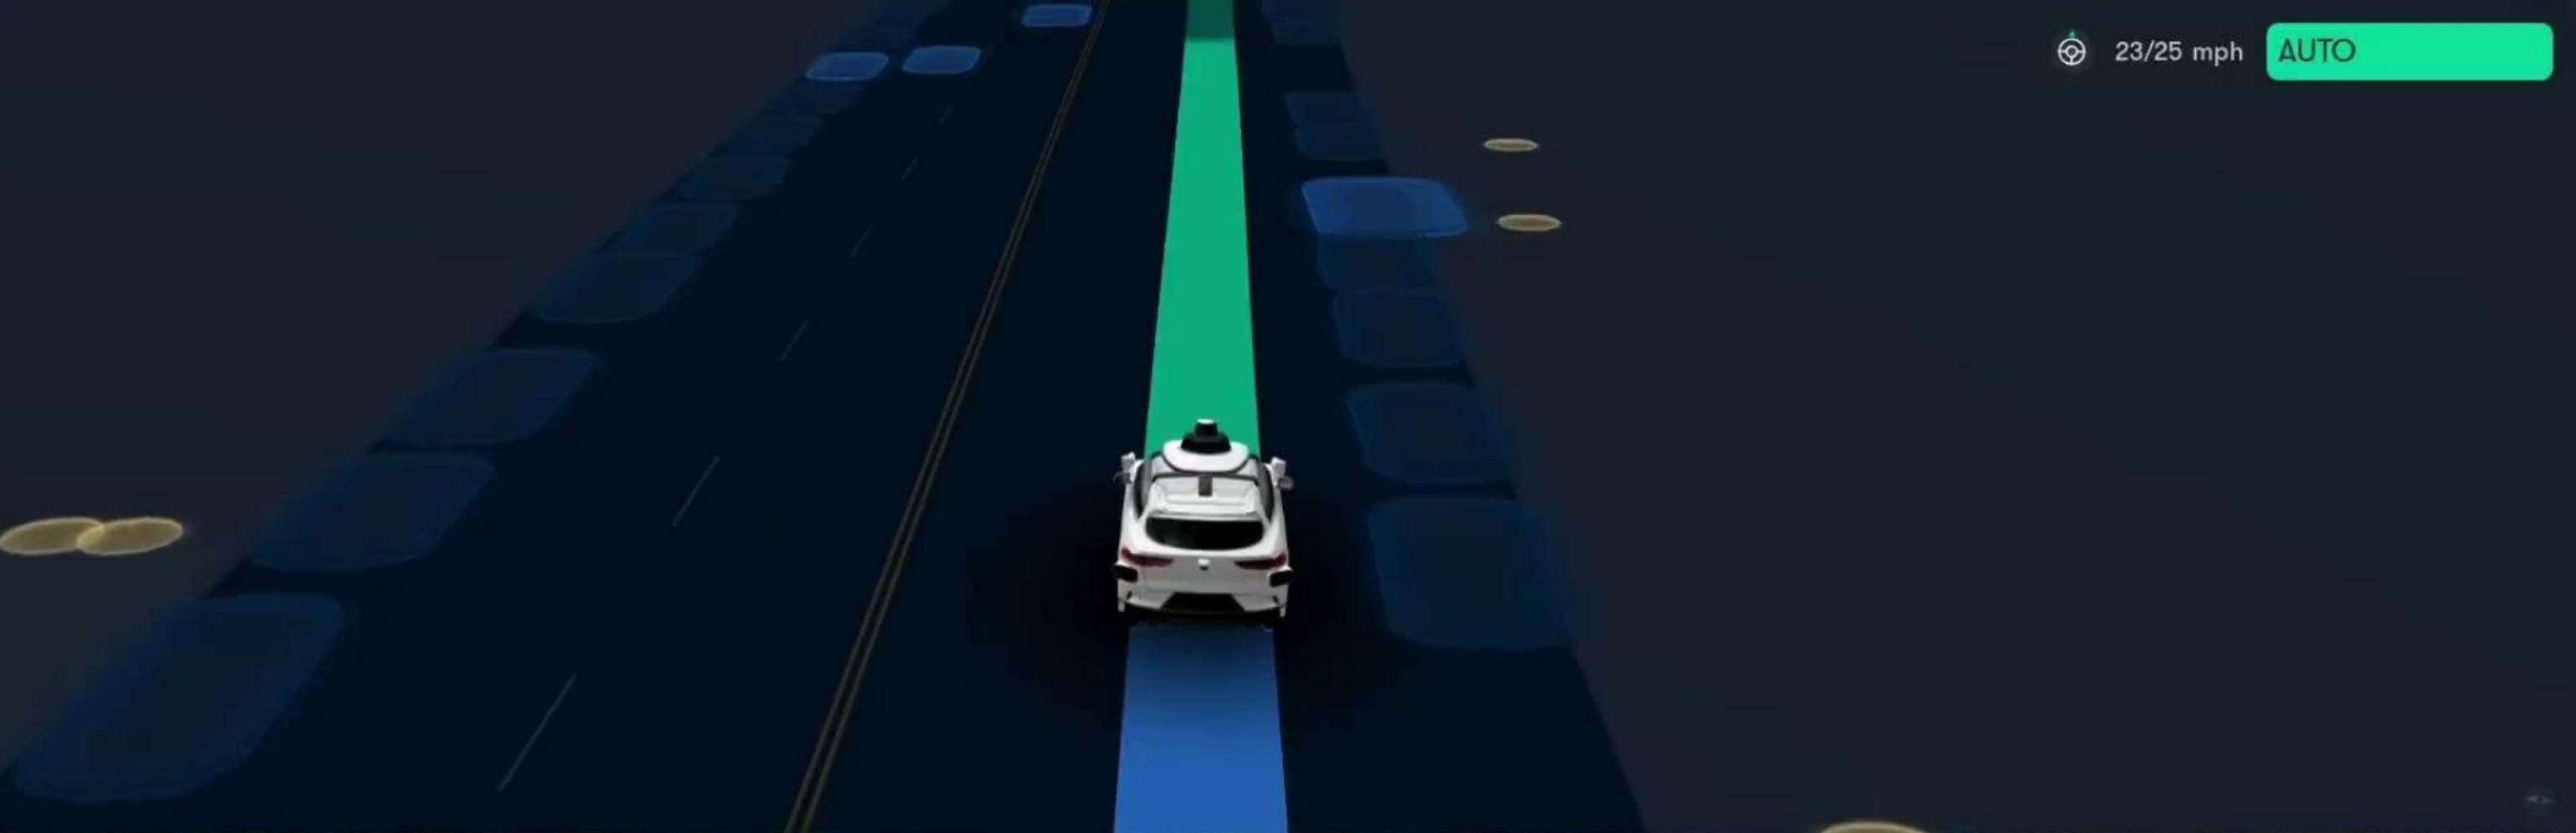
NUSCENES DATASET

Method	self-supervised?	L2 (m) 1s	L2 (m) 2s	L2 (m) 3s	Avg L2 (m)
UniAD [Hu et al., 2023]	✗	0.42	0.64	0.91	0.66
DriveVLM [Tian et al., 2024]	✗	0.18	0.34	0.68	0.40
VAD [Jiang et al., 2023]	✗	0.17	0.34	0.60	0.37
OmniDrive [Wang et al., 2024a]	✗	0.14	0.29	0.55	0.33
DriveVLM-Dual [Tian et al., 2024]	✗	0.15	0.29	0.48	0.31
Ego-MLP* [Zhai et al., 2023]	✓	0.15	0.32	0.59	0.35
BEV-Planner [Li et al., 2024]	✓	0.16	0.32	0.57	0.35
EMMA (random init)	✓	0.15	0.33	0.63	0.37
EMMA	✓	0.14	0.29	0.54	0.32
EMMA+	✓	0.13	0.27	0.48	0.29

Foundation models open up new, scalable avenues for deep semantic understanding of the world and of human behavior.

Ultimately, these advances are in service of building a product that can delight riders and earn their trust every day.

23/25 mph AUTO



If you're interested in building **foundation models** for robots operating in the real world **today**.

Join us on this ride!



vanhoucke@waymo.com



CVPR 2025

Come chat with
me at **Booth 1223!**

June 11-15, 2025

